

Pragmatic Features of Professional Film Criticism

Reid Baumeister

Introduction

Even in today's digital media landscape, professional film criticism continues to shape audience decisions and influence the commercial success of movies. While film criticism is now more accessible than ever through social media and platforms like Rotten Tomatoes, Letterboxd, and Reddit, professional critics provide historical context and editorial independence from the industry.

Of these platforms, Rotten Tomatoes has become one of the most dominant film review websites in the United States, attracting over 100 million monthly visitors. According to Morning Consult, 33% of U.S. adults check Rotten Tomatoes before seeing a film, and 63% report having decided not to see a movie based on its scoring on the site. Additionally, films earning the "Certified Fresh" badge (requiring at least 75% positive reviews from a minimum of 40 professional critics) see an average box office increase of 15-20%.

This influence stems in part from Rotten Tomatoes' key structural feature: professional critics' reviews are aggregated separately from user reviews. This distinction allows audiences to compare expert and amateur perspectives, which often differ not only in evaluation but also in how those evaluations are expressed. While prior research has examined differences in review content, less attention has been given to how professional critics construct their evaluations linguistically. Understanding these patterns has both theoretical value and practical implications, particularly for automating review classification and improving recommendation systems.

This study addresses two research questions. RQ1 (Linguistic Patterns) asks whether critics use different linguistic strategies when writing positive versus negative reviews. Building on prior research that shows professionals use language strategically in domains such as corporate reporting and academic writing (Hyland 2005), this study extends that framework to film criticism. It examines whether critics use hedging and intensification to manage tone and maintain credibility when delivering negative evaluations. RQ2 (Predictive Power) asks whether

these linguistic features can be used to predict whether a review is classified as Fresh or Rotten, and which features are most informative.

Two hypotheses guide this analysis. H1 (Pragmatic Asymmetry) predicts that Rotten reviews will contain more hedging and greater use of intensifiers than Fresh reviews. This expectation is informed by prior work showing that professional writers adjust language in response to evaluative context, often being more indirect when communicating negative information. In academic writing, hedging functions as a form of metadiscourse that allows writers to manage claims and signal appropriate caution (Hyland 2005). This study tests whether similar strategies appear in film criticism.

H2 (Classification Performance) predicts that a supervised machine learning model will achieve greater than 70% accuracy in predicting review valence (Fresh versus Rotten) using pragmatic and syntactic features. This threshold tests whether the linguistic patterns in H1 are observable and robust enough to classify individual reviews. This threshold represents a meaningful improvement over both the 50% random baseline for binary classification and simple sentiment analysis approaches, validating that sophisticated linguistic features add value beyond basic positive-versus-negative word counting. It also provides a benchmark for comparing the effectiveness of linguistic features relative to simpler approaches such as sentiment analysis.

To answer my research questions and test my hypotheses, I used three main text analysis methods (Corpus Linguistics, Machine Learning, Natural Language Processing) on a balanced corpus of 100,000 critic reviews sampled from the Kaggle Rotten Tomatoes dataset.

Literature Review

The concept of hedging as a strategic language choice has been well established in research on professional communication. Hyland's (2005) metadiscourse work explains hedging as a deliberate strategy for writer-reader interactions that serve many functions. Hedging allows writers to build credibility through caution and demonstrate intellectual writing that may be expected in the industry. He distinguishes between epistemic hedging and strategic hedging, where professionals use tentative language to position themselves better- proving that using hedging language can be strategic and shows expertise of a subject.

Complementing this perspective, Li (2008) found that firms with weaker financial results produce less readable and more complex disclosures, suggesting that writers adjust language use strategically in response to negative outcomes. This raises the possibility that similar patterns may emerge in other evaluative contexts, such as film criticism.

Research on film reviews has primarily focused on what critics discuss rather than how they express evaluations. For example, Kim et al. (2023) analyze topic consistency between professional critics and general users, finding that both groups tend to focus on similar themes. This suggests that differences between critics and non-critics may not stem from expertise alone, but potentially from how evaluations are constructed linguistically. As a result, examining the stylistic and pragmatic features of critic reviews offers a complementary perspective to existing content-based analyses.

In addition, industry research highlights the broader impact of review platforms. A report by Morning Consult (2018) shows that Rotten Tomatoes significantly influences viewing decisions, even though its scores do not strongly predict overall box office success. These findings emphasize the importance of understanding how reviews are written, given that numerical scores may obscure linguistic nuance.

Prior research suggests that professional language is both strategic and situational: writers adjust their communication in response to context (Li), using specific linguistic tools to manage stance (Hyland), and produce content that influences audience behavior (Kim). However, important gaps remain. Existing studies focus on academic writing, corporate disclosure, or review content. There is little work examining how critics use pragmatic features such as hedging, intensification, and contrast when expressing positive versus negative judgments. Prior research has not tested whether these linguistic patterns are strong enough to support automated classification at the individual review level. This study addresses these gaps by applying linguistic theory and computational methods to a large corpus of Rotten Tomatoes reviews. By analyzing review features, it extends existing research into a new domain and evaluates whether patterns in professional criticism are measurable and useful.

Data & Methods

This study uses the Kaggle Rotten Tomatoes dataset, a publicly available corpus containing over 1.13 million critic reviews published between 2000 and 2020. Each record includes a review excerpt, Fresh/Rotten classification, critic name, publication source, and review date. From this dataset, a balanced subset of 100,000 reviews was created, consisting of 50,000 Fresh and 50,000 Rotten reviews. Additional preprocessing steps included filtering for reviews with a minimum length of 150 characters and limiting each publication to a maximum of 500 reviews to reduce source bias.

A key constraint of the dataset is the short length of Rotten Tomatoes excerpts, with an average of 21 words per review. This creates substantial feature sparsity, particularly for low-frequency linguistic features such as hedging, with roughly 50% of reviews containing no hedge terms. This reflects a structural characteristic of the platform rather than a limitation of the dataset, but it has important implications for feature detection and model performance.

Using the spaCy natural language processing library, 11 pragmatic and syntactic features were extracted from each review. These include hedge density, verb density, additive and contrastive markers, and counts for strong and weak intensifiers. These features create a representation of each review's linguistic characteristics, allowing the classifier to distinguish reviews based on how critics write rather than simply which specific words they choose.

To address the research questions, three complementary analytical approaches were employed. First, corpus linguistics methods were used to examine linguistic patterns and test H1 (Pragmatic Asymmetry). This included frequency analysis, keyword analysis, and collocation analysis. Frequency analysis compares feature densities across Fresh and Rotten reviews, focusing on hedging and intensification. Keyword analysis identifies words that are statistically overrepresented in each category, while collocation analysis examines how hedging terms function in context by analyzing co-occurring words.

Second, a Support Vector Machine (SVM) classifier was used to test H2 (Classification Performance). The model was trained on the extracted pragmatic and syntactic features to predict whether a review is Fresh or Rotten. An 80/20 train-test split was applied, maintaining class


balance across both sets. SVM allows for interpretation of feature importance and assessment of whether linguistic style alone is sufficient for classification.

Third, sentiment analysis was conducted using the VADER model to provide a baseline comparison, classifying reviews based on sentiment rather than pragmatic features. Comparing sentiment-based performance with the SVM results allows for evaluation of whether more complex linguistic features provide additional predictive value beyond simple polarity detection.

Results

Method 1: Corpus Linguistics


1.1: Frequency Comparison



The frequency analysis compared hedging and intensifier usage across Fresh and Rotten reviews per 100 words. Rotten reviews contained 19.2% more hedging language than Fresh reviews, which is lower than the hypothesized 50% increase. Similarly, Fresh reviews contained only 1.06 times more strong intensifiers than Rotten reviews, compared to the predicted twofold difference.

Overall, the results partially support H1, with effects in the expected direction but smaller than anticipated.


1.2 Keyword Analysis



Keyword analysis using the log-likelihood ratio (G^2) identified words that were disproportionately associated with each review category. Fresh reviews were characterized by terms such as “best,” “entertaining,” “performance,” and “great,” suggesting a focus on praise and recognition of specific film elements such as acting and overall execution. The prominence of terms like “performance” and “portrait” suggests that critics frequently respond positively to films emphasizing character development and emotional depth. In contrast, Rotten reviews were associated with terms such as “fails,” “bad,” “dull,” “boring,” and “worst,” indicating direct negative evaluation. Words such as “unfortunately” and “too” suggest that negative reviews are often framed as critiques of execution rather than absolute dismissal, implying disappointment relative to expectations rather than outright rejection.


Together, these patterns suggest that Fresh reviews tend to emphasize strengths and positive attributes of film components, while Rotten reviews highlight unmet expectations.

1.3: Collocation Analysis



Collocation analysis examined how critics use hedging terms when evaluating films in their reviews. Across the dataset, a consistent pattern emerged in which hedging expressions frequently co-occurred with contrastive markers, particularly “but.” Rather than solely indicating uncertainty, hedges are used to introduce contrast between positive and negative assessments within a review. This pattern suggests that critics rarely present evaluations as purely positive or negative, but attempt to balance praise with criticism rather than making absolute claims. The surrounding collocates further support this interpretation. Across multiple hedge terms, the repeated presence of contrastive markers indicates that professional film criticism relies heavily on qualified assessment rather than blunt judgment. Critics use hedging to frame transitions between competing evaluations, allowing them to recognize both strengths and weaknesses within the same review. Overall, these findings suggest that hedging in film criticism is a strategy for managing evaluative balance, rather than simply signaling uncertainty. Professional film criticism is more than opinion; it is a persuasive analysis with carefully qualified judgments.

Method 2: Machine Learning



The SVM model achieved an overall accuracy of 56%, with a precision of 0.564, recall of 0.530, and an F1 score of 0.546. This performance is only slightly better than random chance, suggesting that writing style is a weak predictor of review sentiment.

Content and specific evaluative language are more valuable for determining review outcome. Additive markers are by far the strongest predictors of Fresh reviews; critics build positive arguments by stacking multiple praises rather than relying on one dominant point. Especially in these shorter reviews, critics want to mention several aspects, which can help differentiate the film and highlight their own insights. Strong intensifiers and adjective density are other positive predictors. Fresh reviews tend to be more expressive and evaluative, using stronger descriptive language and emphasis.

Features associated with rotten reviews include adverb, verb, and hedging density, as well as contrastive markers and weak intensifiers. Rotten reviews are more likely to explain, contextualize, and contrast, supporting the idea that negative reviews are less direct and more

linguistically complex. While hedging and writing style contribute to review sentiment, it is not enough on its own to reliably classify reviews.

Method 3: Natural Language Processing

The results show that fresh reviews have an average sentiment score of +0.342, while Rotten reviews average -0.189. The difference between these two groups is statistically significant, meaning separation is highly unlikely to be due to chance. The model produced a classification accuracy of 64.2%, outperforming the SVM model. Sentiment is a stronger predictor of review outcomes than general linguistic style in this short-form context. Pragmatic features capture more subtle variation in how critics write but contribute less predictive power at the individual level.

Discussion and Business Implications

The results provide partial support for H1 (Pragmatic Asymmetry). Rotten reviews contained 19% more hedging language than Fresh reviews, which is directionally consistent with my hypothesis but substantially smaller than the predicted 50% effect. Similarly, Fresh reviews showed only a modest increase in strong intensifiers (1.06x). While these findings suggest that critics adjust linguistic strategies based on evaluation, the magnitude of the effect is limited, indicating that pragmatic differences exist but are relatively subtle in this dataset.

H2 (Classification Performance) is not supported by the machine learning results. The SVM model achieved 56% accuracy, falling short of the 70% threshold and barely above random chance. Sentiment analysis performed better with 64.2% accuracy. This suggests that in extremely short texts, sentiment provides a stronger signal for classification than structural or pragmatic features. While pragmatic features capture variation in how critics construct arguments, they are less effective for prediction at the individual review level due to limited text length and feature sparsity.

Despite this, collocation analysis reveals a consistent pattern in how hedging is used. The frequent co-occurrence of hedging with contrastive markers such as “but” suggests that hedging functions less as an expression of uncertainty and more as a rhetorical device for structuring balanced evaluation. Critics appear to use hedging to acknowledge multiple dimensions of a

film, integrating both positive and negative assessments within a single coherent argument. This reinforces prior research on strategic language use in professional contexts, including work by Ken Hyland, as well as the broader idea that professional communication is shaped by evaluative constraints.

These findings have practical implications for review aggregation platforms such as Rotten Tomatoes. Sentiment models provide a more efficient baseline for automating review classification in short-text environments. Pragmatic features may be more valuable as interpretive signals rather than predictive inputs, helping identify nuanced or mixed reviews.

These results highlight a structural constraint of modern review ecosystems: as platforms increasingly favor shorter reviews (e.g., social media and user-generated content platforms), sophisticated linguistic signals become harder to detect and less useful for prediction. This suggests that the value of linguistic complexity depends strongly on text length, with longer-form reviews likely offering richer opportunities for pragmatic analysis.

Conclusion

This study examined how professional film critics use language, focusing on hedging and intensifiers in reviews and whether these patterns can support automated classification. Using a balanced dataset of 100,000 Rotten Tomatoes critic reviews, the results show that critics adjust their writing based on evaluation, but not as strongly as expected. Rotten reviews contained 19% more hedge words than Fresh reviews. This result was smaller than expected, possibly due to the small character count of each review. The collocation analysis revealed how critics rely on hedging to structure their arguments, often acknowledging strengths before delivering criticism. Hedging is a strategic action to balance evaluation, not a sign of hesitation. While these differences are strategically significant, they are small in absolute terms, with only about 0.04 more hedges per review. As a result, the SVM model performed slightly better than chance while VADER sentiment analysis improved to 64%. For short texts, a simple sentiment analysis performs better and is more efficient for classification. Overall, this study extends research on strategic language use into film criticism, highlights linguistic trends for professional critics, and offers insights for platforms like Rotten Tomatoes.

There are several limitations to consider. The short length of reviews limits linguistic variation and changes the way we must interpret the results. The analysis also only focuses on a single platform. I believe that the results would differ if taken from sites like Letterboxd where the reviews are mainly from users and are written differently. Future research could examine longer reviews, compare critics and general users directly, and explore how these patterns may have changed over time.

The findings in this paper suggest that professionalism in evaluative writing is focused on how language is used, often through hedges and intensifiers. Critics write differently depending on their rating of the film, and Rotten Tomatoes can adjust their recommendation algorithms to better showcase impactful reviews. This has implications for how review platforms interpret and surface critic content, especially as digital media continues to favor shorter-form writing.

References

Hyland, K. (2005). *Metadiscourse: Exploring interaction in writing*. Continuum.

Kim, E., Ding, M., Wang, X., & Lu, S. (2023). Does topic consistency matter? A study of critic and user reviews in the movie industry. *Journal of Marketing*, 87(3), 428–450.

Li, F. (2008). Annual report readability, current earnings, and earnings persistence. *Journal of Accounting and Economics*, 45(2–3), 221–247.

Morning Consult. (2018). America can stomach Rotten Tomatoes. *Morning Consult Pro*.
<https://pro.morningconsult.com/articles/america-can-stomach-rotten-tomatoes>

Rotten Tomatoes. (n.d.). About: The Tomatometer rating system.
<https://www.rottentomatoes.com/about>